

Databases, SQL, and

cs109, Fall 2015 (#cs109)

Rahul Dave

`rahuldave@gmail.com`, `@rahuldave`, `staff@cs109.org`

ANNOUNCEMENTS

Class in in Science Center B starting THIS thursday, 17th Sep, 2015!

It took about three years before the BellKor's Pragmatic Chaos team managed to win the prize ... The winning algorithm was ... so complex that it was never implemented by Netflix.¹

¹ <https://hbr.org/2012/10/big-data-hype-and-reality>

Machine

Data Management

Data Mining

Machine Learning

Visualization

Business Intelligence

Statistics

Human

Human Cognition

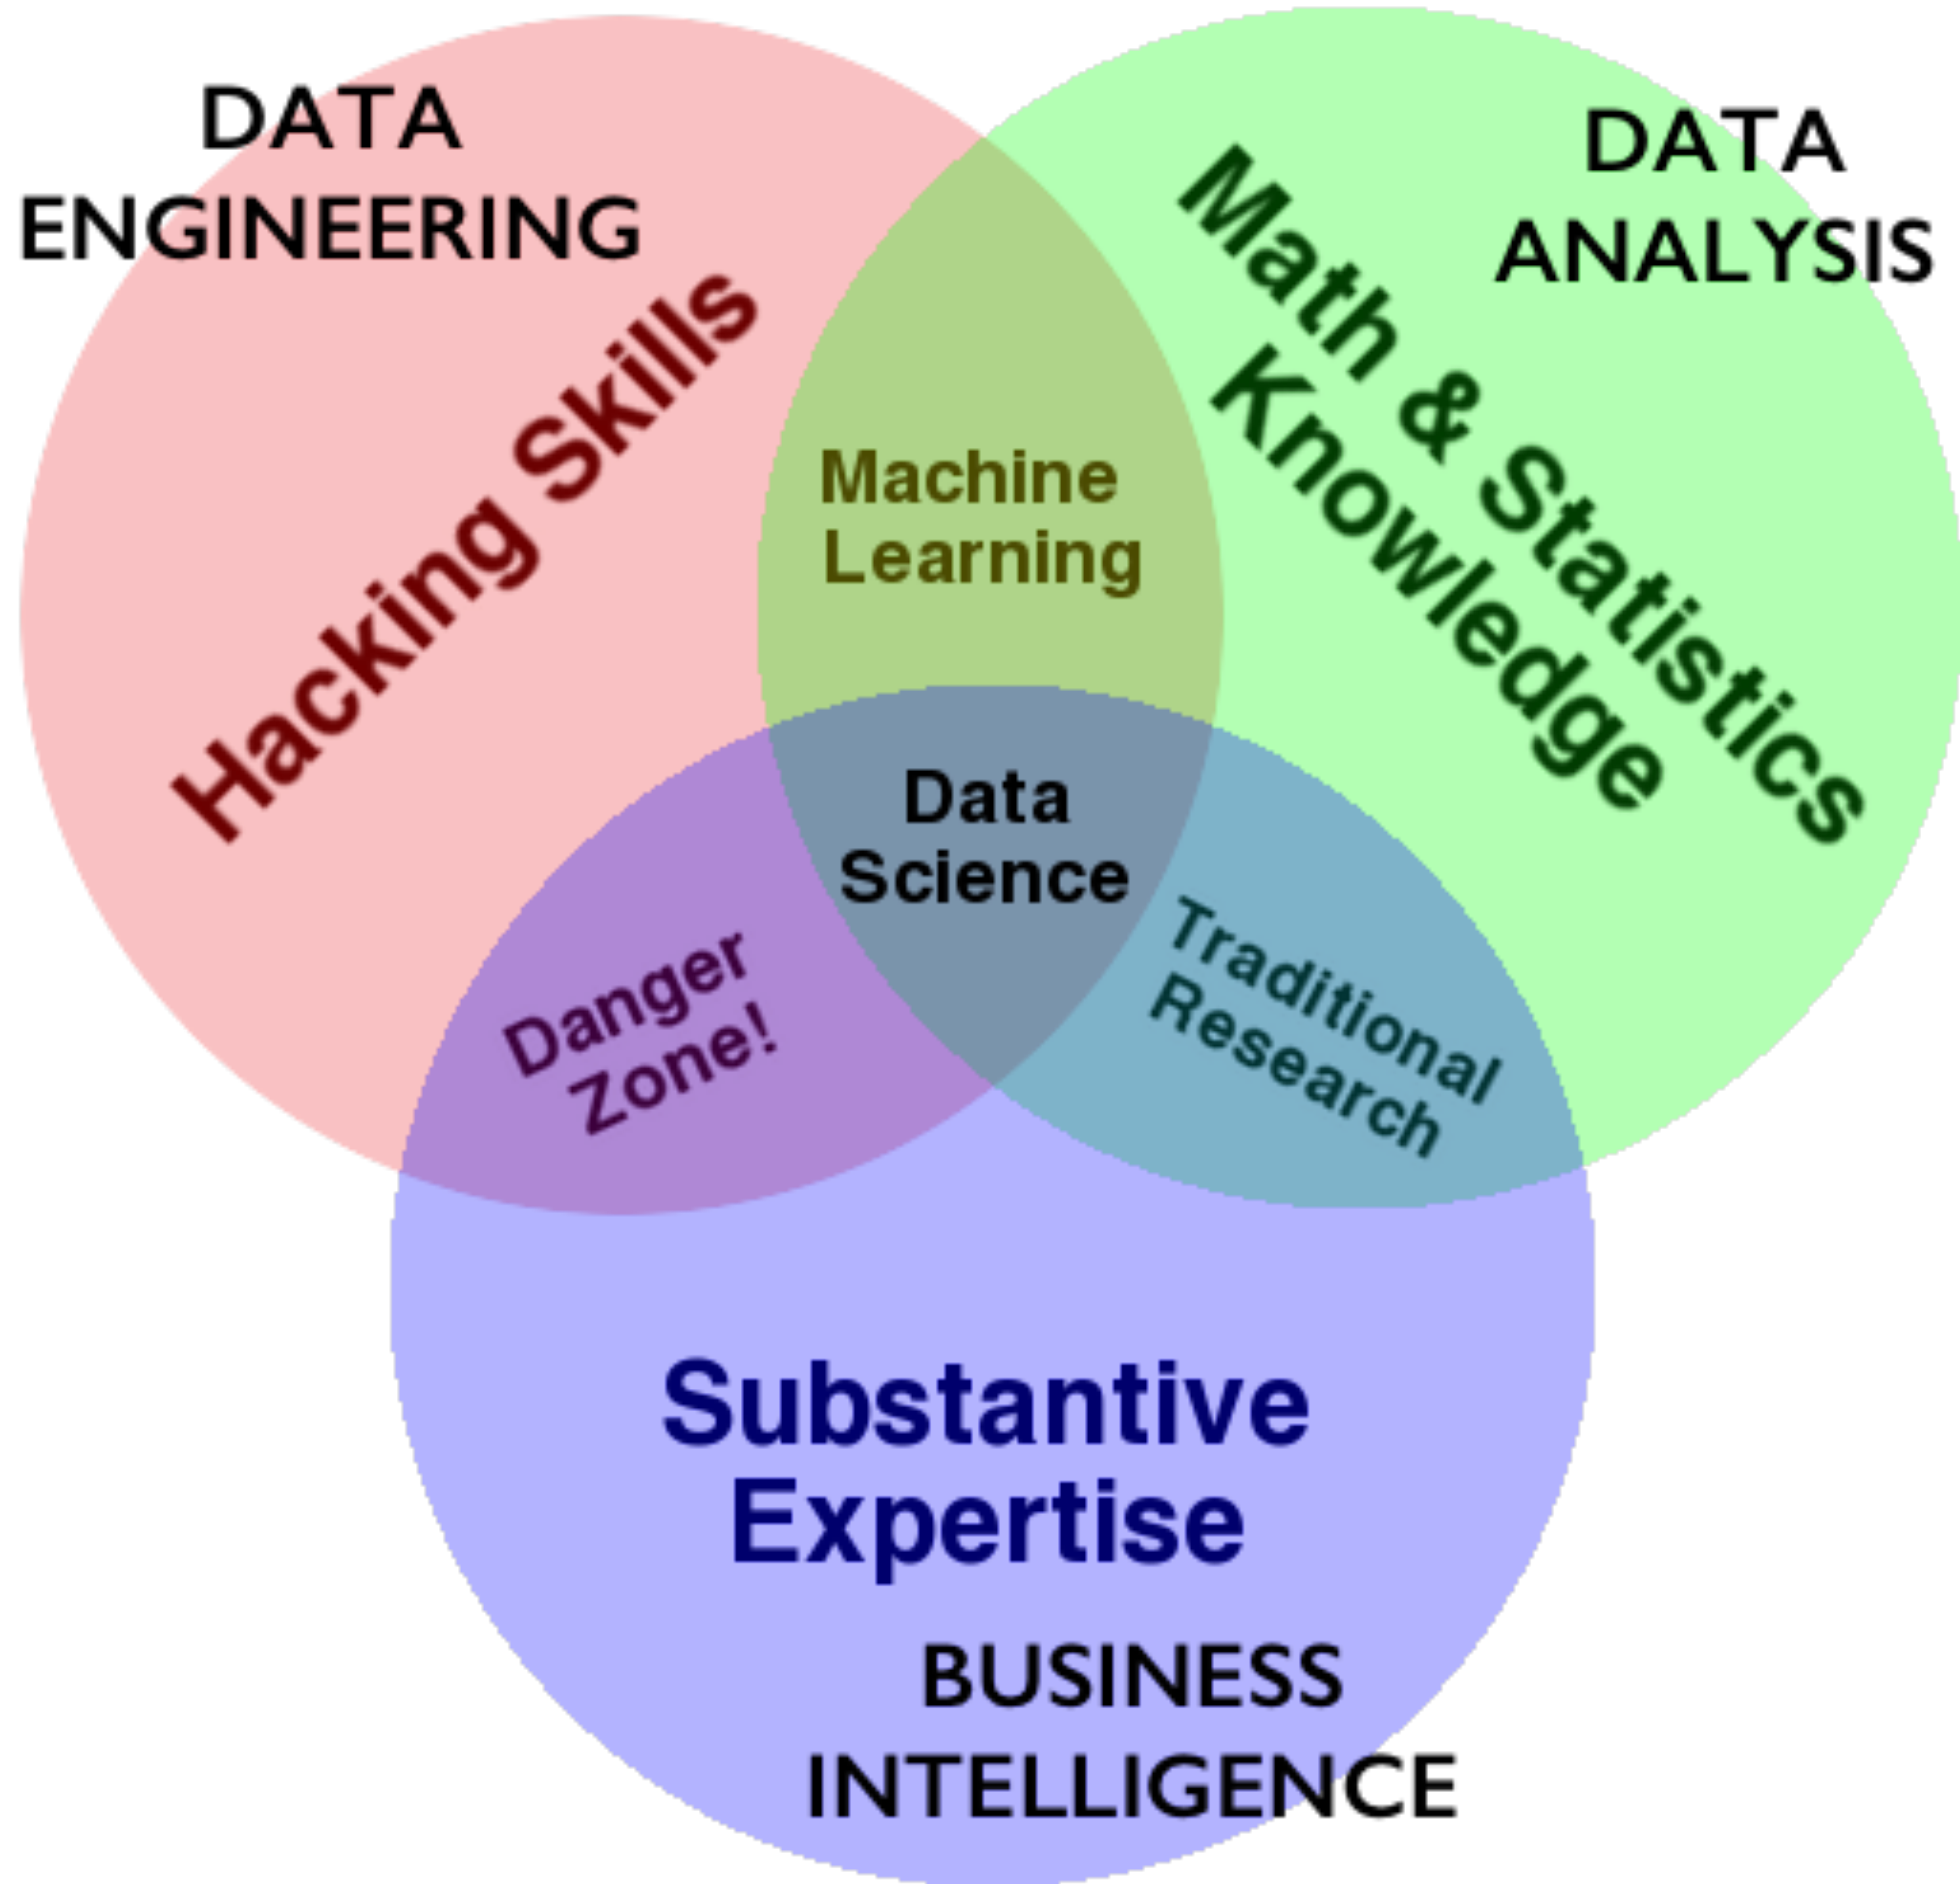
Perception

Story Telling

Decision Making
Theory

Data Science





Data Scientist: Sexiest Job of the 21st Century

It's important that our data team wasn't comprised solely of mathematicians and other "data people." It's a fully integrated product group that includes people working in design, web development, engineering, product marketing, and operations. They all understand and work with data, and I consider them all data scientists... Often, an engineer can have the insight that makes it >clear how the product's design should work, or vice-versa — a designer can have the insight that helps the engineers understand how to better use the data. Or it may take someone from marketing to understand what a customer really wants to accomplish.²

² D. J. Patil, U.S. Chief Data Scientist, Building data science teams. " O'Reilly Media, Inc.", 2011.

DATA ENGINEERING

- **compute:** code, python, R, julia, spark, hadoop
- **storage/database:** git, SQL, NoSQL, HBase, disk, memory
- **devops:** AWS, docker, mesos, repeatability
- **product:** database, web, API, viz, UI, story

Different at different scales....

What kind of data storage do you need?

- **memory**
- **disk:** what if we do not fit?
- **cluster:** what if we still do not fit?
- **cluster:** what if we need/can use parts?
- What if we **MUST** bring compute to disk?

What kind of data access do you need?

- **relational:** pandas, SQL: Postgres, sqlite, Hbase, VoltDB
- **document oriented:** MongoDB, CouchDB
- **key-value:** Riak, Redis, Memcached
- **graph oriented:** Neo4J

Today we'll focus on relational

- What is a relational Database?
- What Grammar of Data does it follow?
- How is this grammar implemented in Pandas?
- How is this grammar implemented in SQL

Relational Database

Dont say: seek 20 bytes onto disk and pick up from there. The next row is 50 bytes hence

Say: select data from a set. I dont care where it is, just get the row to me.

Relational Database(contd)

- A collection of tables related to each other through common data values.
- Rows represent attributes of something
- Everything in a column is values of *one* attributes
- A cell is expected to be atomic
- Tables are related to each other if they have columns called keys which represent the same values

Contributors

Table:

		id	first_name	last_name	middle_name	party
Filter		Filter	Filter	Filter	Filter	
1	1	16	Mike	Huckabee		R
2	5	20	Barack	Obama		D
3	6	22	Rudolph	Giuliani		R
4	7	24	Mike	Gravel		D
5	8	26	John	Edwards		D
6	9	29	Bill	Richardson		D
7	10	30	Duncan	Hunter		R
8	11	31	Dennis	Kucinich		D
9	12	32	Ron	Paul		R

Candidates

Scales of Measurement

- Quantitative (Interval and Ratio)
- Ordinal
- Nominal³

Scale	Basic Empirical Operations	Mathematical Group Structure	Permissible Statistics (invariantive)
NOMINAL	Determination of equality	Permutation group $x' = f(x)$ $f(x)$ means any one-to-one substitution	Number of cases Mode Contingency correlation
ORDINAL	Determination of greater or less	Isotonic group $x' = f(x)$ $f(x)$ means any monotonic increasing function	Median Percentiles
INTERVAL	Determination of equality of intervals or differences	General linear group $x' = ax + b$	Mean Standard deviation Rank-order correlation Product-moment correlation
RATIO	Determination of equality of ratios	Similarity group $x' = ax$	Coefficient of variation

³ S. S. Stevens, Science, New Series, Vol. 103, No. 2684 (Jun. 7, 1946), pp. 677-680

Grammar of Data

Been there for a while (SQL, Pandas), formalized in `dplyr`⁴.

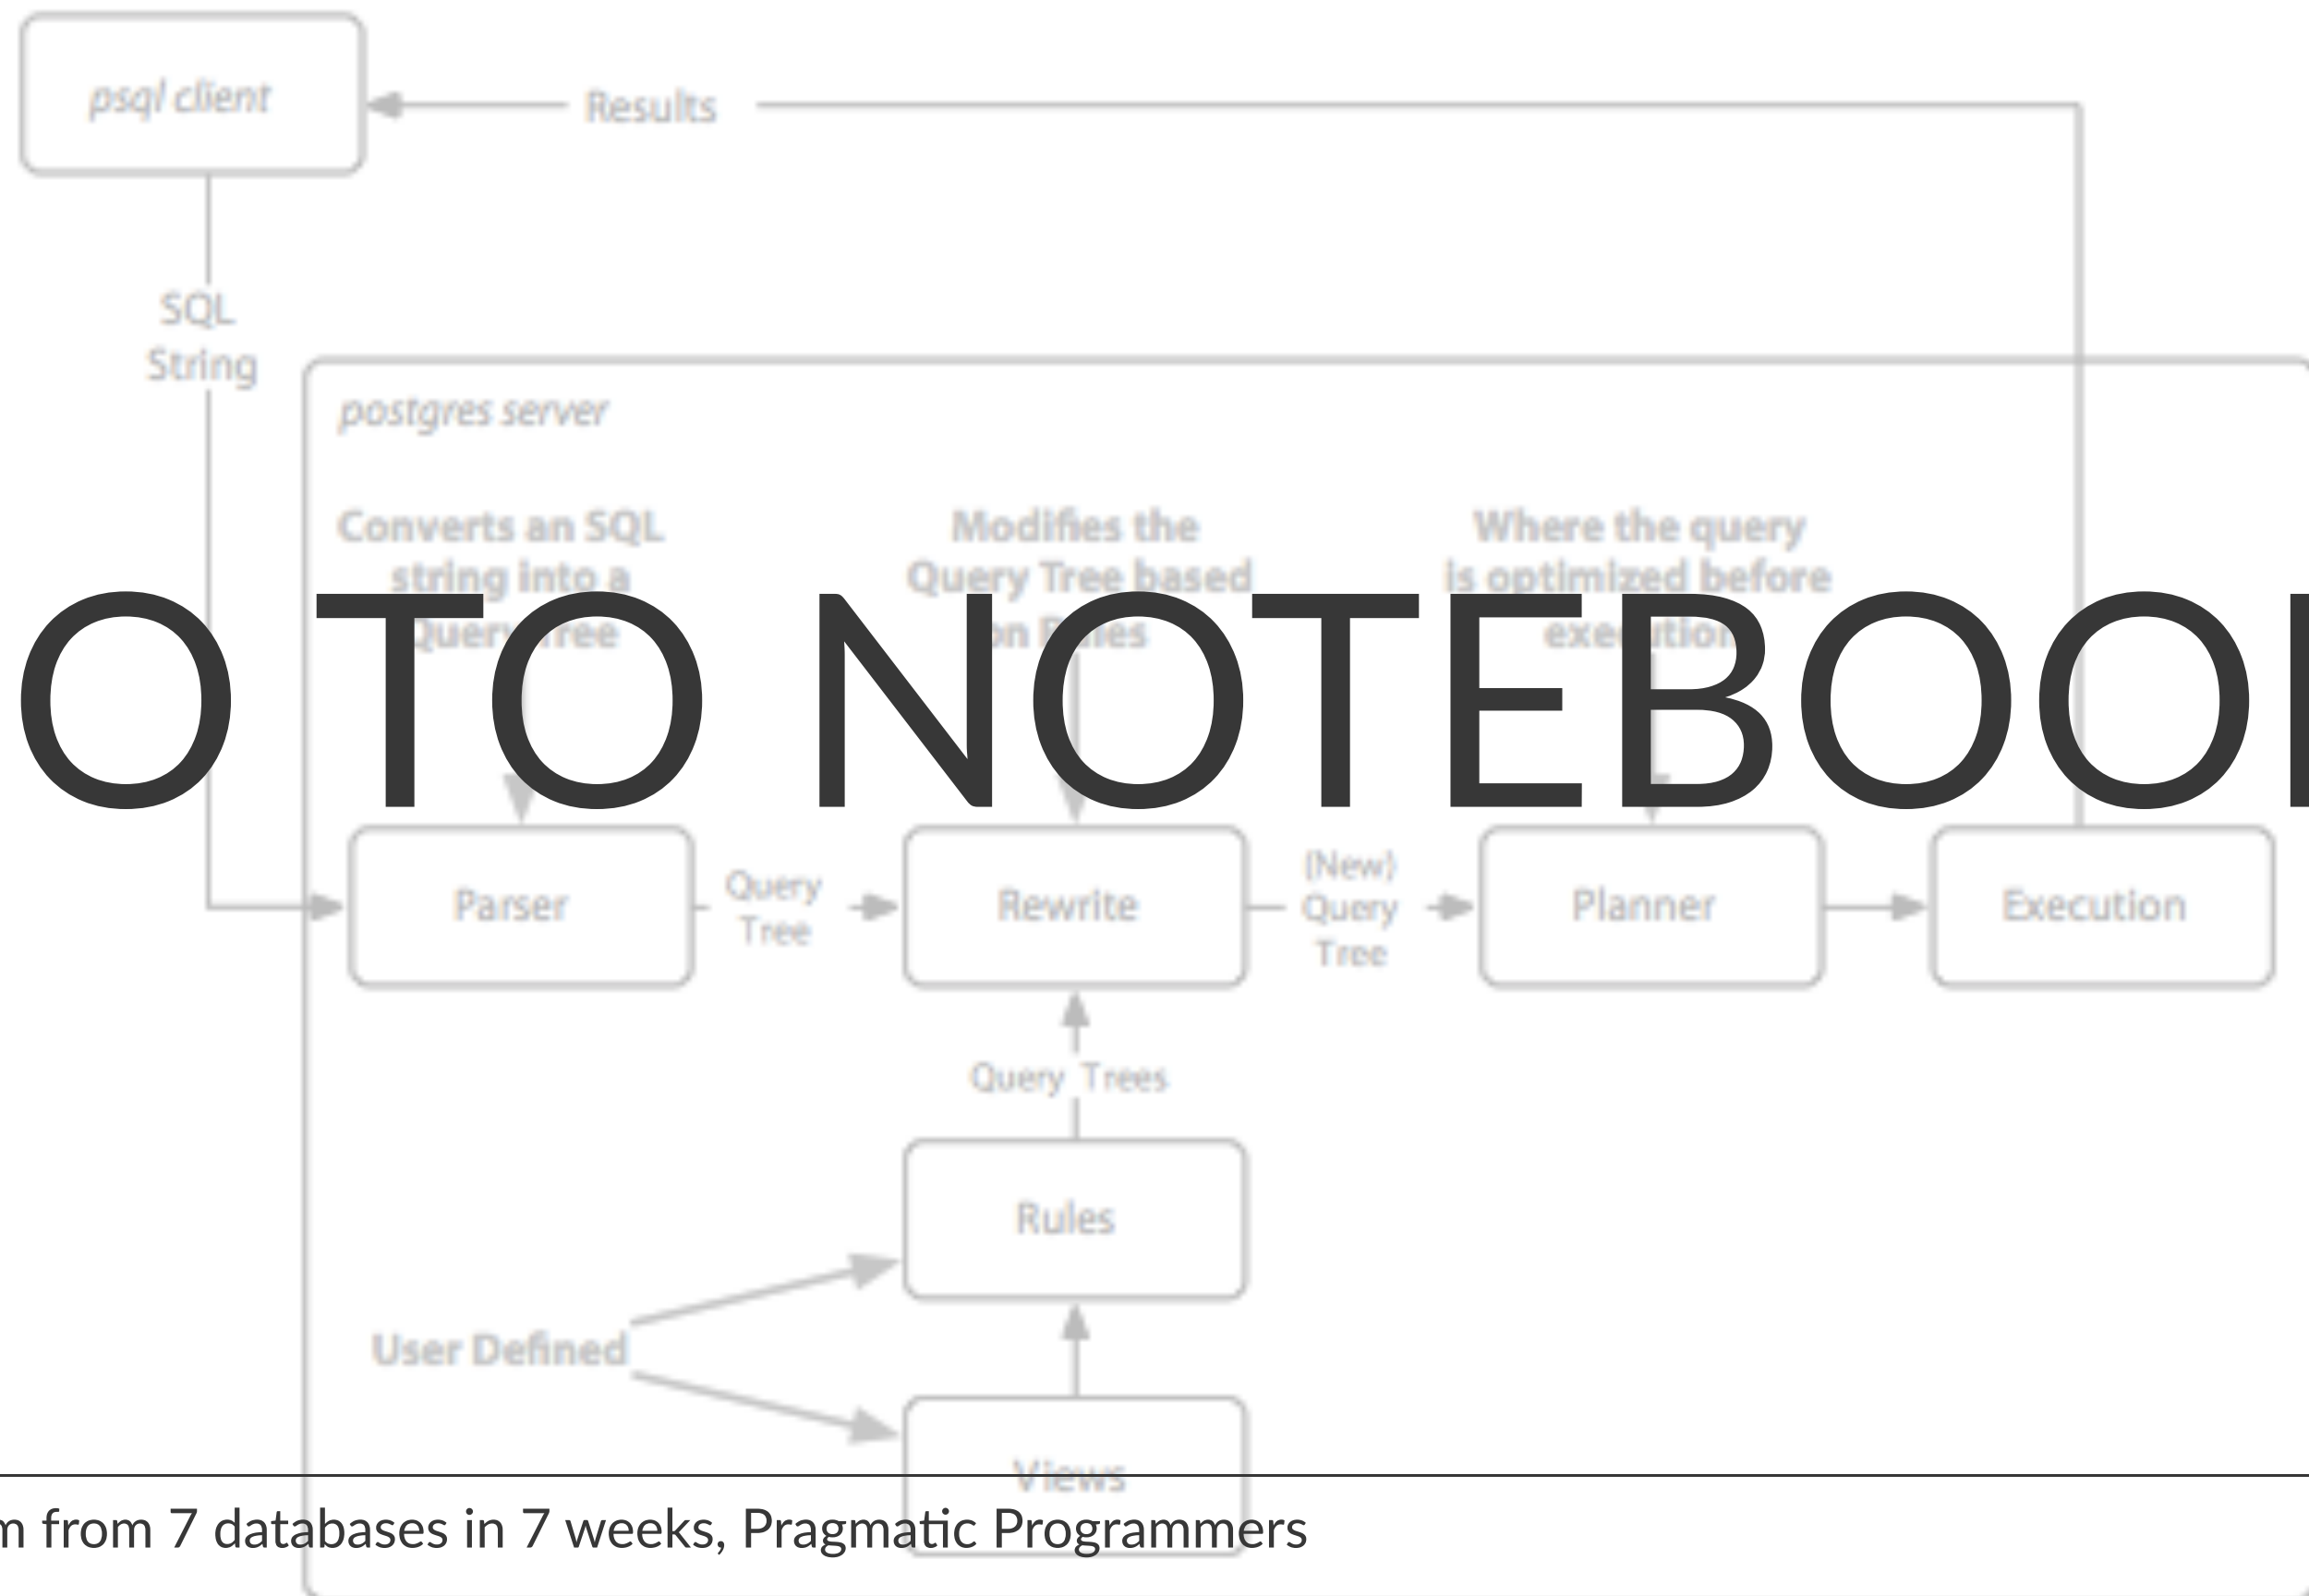
- provide simple verbs for simple things. These are functions corresponding to common data manipulation tasks
- second idea is that backend does not matter. Here we constrain ourselves to Pandas and sqlite
- multiple backends implemented in Pandas, Spark, Impala, Pig, dplyr, ibis, blaze

⁴ Hadley Wickham: <https://cran.rstudio.com/web/packages/dplyr/vignettes/introduction.html>

Why bother

- learn how to do core data manipulations, no matter what the system
- relational databases critical for non-memory fits. Big installed base.
- one off questions: google, stack-overflow, <http://chrisalbon.com>

GO TO NOTEBOOK⁵



⁵ Diagram from 7 databases in 7 weeks, Pragmatic Programmers

RDBMS when:

- data structure regularity is known
- transactions are required
- benefit from years of tuning
- not good for deep hierarchy
- which kind depends on use case: pandas, hbase, columnar, postgres,...

F

I

N